

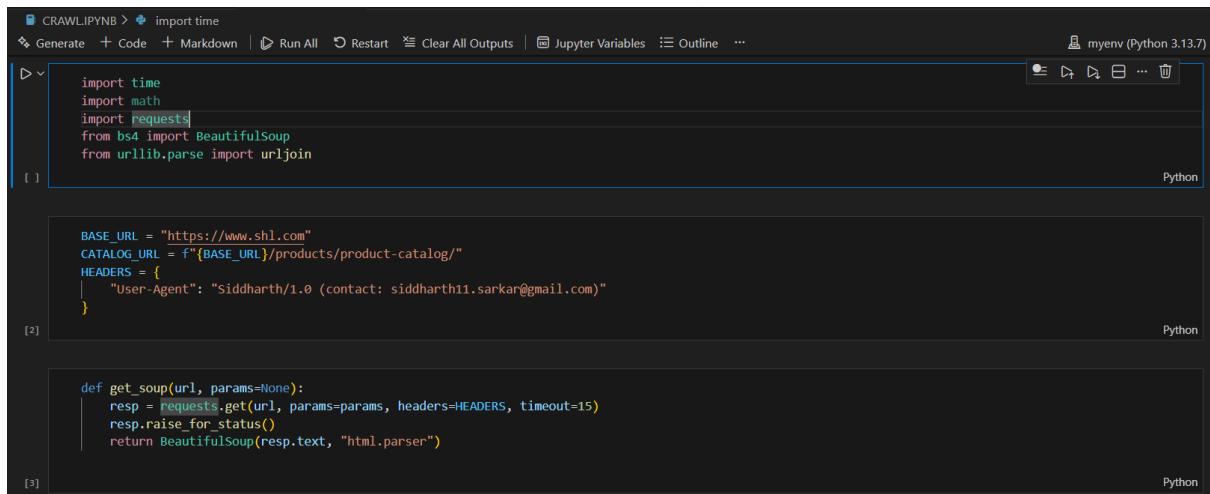
# SHL ASSESSMENT

## RESEARCH AI INTERN

### SIDDHARTH SARKAR

Before building the rag pipeline we need to crawl the data for individual test solutions from shl's product catalog

Sample code snippet:



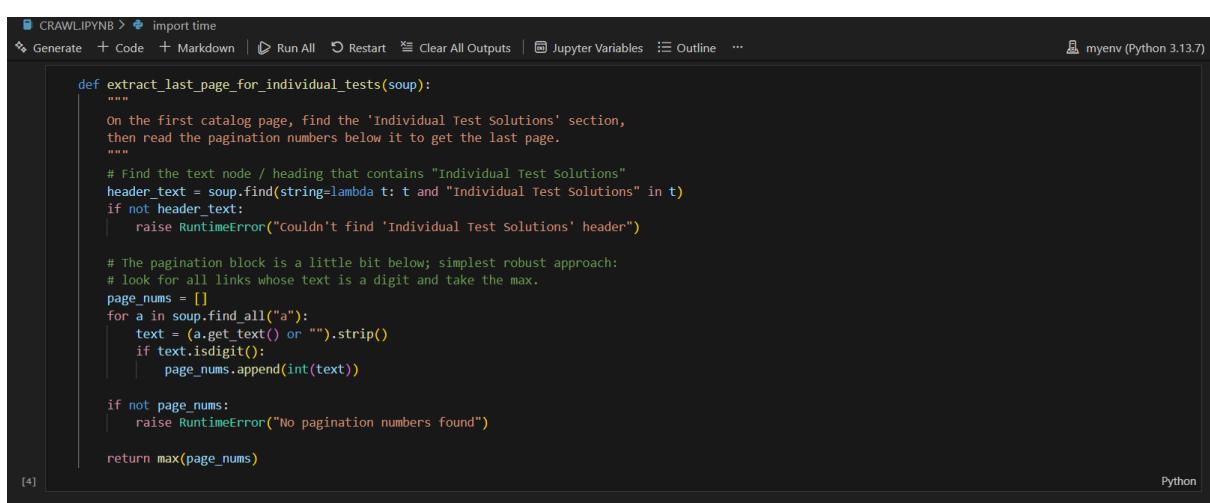
The screenshot shows a Jupyter Notebook interface with two code cells. The first cell contains code to import modules and define constants for the base URL and catalog URL. It also defines a header dictionary with a user-agent string. The second cell contains a function to get a BeautifulSoup object from a given URL. The notebook is titled 'CRAWLIPYNB' and is running in a 'myenv (Python 3.13.7)' environment.

```
import time
import math
import requests
from bs4 import BeautifulSoup
from urllib.parse import urljoin

BASE_URL = "https://www.shl.com"
CATALOG_URL = f"{BASE_URL}/products/product-catalog/"
HEADERS = {
    "User-Agent": "Siddharth/1.0 (contact: siddharth11.sarkar@gmail.com)"
}

def get_soup(url, params=None):
    resp = requests.get(url, params=params, headers=HEADERS, timeout=15)
    resp.raise_for_status()
    return BeautifulSoup(resp.text, "html.parser")
```

[2]



The screenshot shows a third code cell in the Jupyter Notebook. This cell contains a function named 'extract\_last\_page\_for\_individual\_tests' which takes a BeautifulSoup object 'soup' as input. The function finds the 'Individual Test Solutions' section and reads the pagination numbers below it to determine the last page. It then looks for all links whose text is a digit and takes the maximum value. If no page numbers are found, it raises a RuntimeError. The notebook is titled 'CRAWLIPYNB' and is running in a 'myenv (Python 3.13.7)' environment.

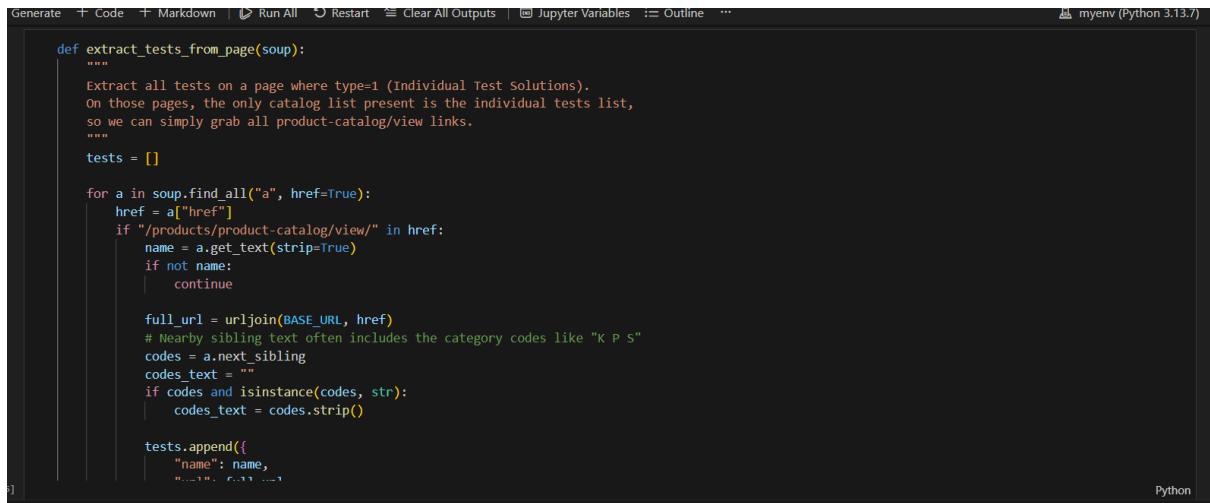
```
def extract_last_page_for_individual_tests(soup):
    """
    On the first catalog page, find the 'Individual Test Solutions' section,
    then read the pagination numbers below it to get the last page.
    """
    # Find the text node / heading that contains "Individual Test Solutions"
    header_text = soup.find(string=lambda t: t and "Individual Test Solutions" in t)
    if not header_text:
        raise RuntimeError("Couldn't find 'Individual Test solutions' header")

    # The pagination block is a little bit below; simplest robust approach:
    # look for all links whose text is a digit and take the max.
    page_nums = []
    for a in soup.find_all("a"):
        text = (a.get_text() or "").strip()
        if text.isdigit():
            page_nums.append(int(text))

    if not page_nums:
        raise RuntimeError("No pagination numbers found")

    return max(page_nums)
```

[4]



```
def extract_tests_from_page(soup):
    """
    Extract all tests on a page where type=1 (Individual Test Solutions).
    On those pages, the only catalog list present is the individual tests list,
    so we can simply grab all product-catalog/view links.
    """
    tests = []

    for a in soup.find_all("a", href=True):
        href = a["href"]
        if "/products/product-catalog/view/" in href:
            name = a.get_text(strip=True)
            if not name:
                continue

            full_url = urljoin(BASE_URL, href)
            # Nearby sibling text often includes the category codes like "K P S"
            codes = a.next_sibling
            codes_text = ""
            if codes and isinstance(codes, str):
                codes_text = codes.strip()

            tests.append({
                "name": name,
                "url": full_url,
                "category_codes": codes_text
            })

    return tests
```

## 1. Problem Understanding

The objective of the project was to develop an **AI-powered recommendation system** capable of retrieving the most relevant **assessment URLs** based on a user query. The solution had to:

- Understand requirements that span multiple skill domains (technical + behavioral + managerial)
- Retrieve and recommend diverse and relevant assessments
- Ensure **non-repetitive results**
- Support **evaluation and performance benchmarking** using **Mean Recall@K**
- Improve accuracy iteratively based on the **Train-Set dataset**

The challenge was ensuring that the system does not hallucinate or recommend repeated content but instead produces **meaningful, diverse recommendations** aligned with expected results

## 2. Solution Architecture

### 2.1 Retrieval-Augmented Generation (RAG) Pipeline

The final pipeline included:

Component	Description
Document Loaders	Loaded catalog assessments with names + URLs
Embedding Model	SentenceTransformers all-mpnet-base-v2
Vector Database	Pinecone (Cosine similarity metric)
Retriever	Max Marginal Relevance (MMR) for diversity
LLM	Hugging Face model integrated with LangChain
Ranking Logic	Return top-K distinct URLs (unique enforced)

### Pipeline Flow

1. Query input
2. Embedding → vector search in Pinecone
3. Retrieval of diverse relevant assessments using **MMR**
4. LLM reranking and structured formatting enforcement
5. Output: Top K recommendations with **unique URLs**

### 3. Performance Evaluation Strategy

#### Metric Selected: Mean Recall@K

Since the goal was to retrieve **all expected relevant assessments**, Recall@K was the correct primary metric.

$$Recall@K = \frac{\text{Relevant items in Top K}}{\text{Total Relevant Items}}$$

$$MeanRecall@K = \frac{1}{N} \sum_{i=1}^N Recall@K_i$$

### Evaluation Setup

- **TrainData1.xlsx** → stored predicted URLs (1 URL per row)
- **Gen\_AI Dataset.xlsx** → stored ground-truth relevant URLs (1 row per URL)
- K = 10

- Normalized all URLs (lowercase, remove query strings, remove slashes) to avoid false mismatches

## 5. Optimization Efforts

Problem Identified	Optimization Applied	Result Improvement
Duplicate documents in index	Stopped re-upserting after first load ( <code>from_existing_index</code> )	Cleaner ranking & better recall
Returned repeated URLs	Enforced unique URLs through LLM prompt + post-processing	5 distinct recommendations
Similarity search lacked diversity	Switched to <b>MMR retriever</b> ( <code>search_type="mmr"</code> )	More coverage across domain
Evaluation false negatives due to formatting	Added URL normalization	Correctly counted matches
Top results missed relevant items	Increased <code>fetch_k</code> and tuned <code>lambda_mult</code>	More exploration depth